

# PROJECT - DEEP LEARNING PROJECT FASHION MNIST CLASSIFICATION.

**STUDENT NAME :** CHINTAREDDY VARSHITHA

**Title:** Deep Learning Project: Fashion MNIST Classification

## **Abstract:**

The Fashion-MNIST dataset is a collection of 70,000 grayscale images of 10 different fashion items. The goal of this project is to build a deep learning model that can classify these images with high accuracy.

## **Objective:**

The objective of this project is to build a deep learning model that can classify Fashion-MNIST images with an accuracy of at least 90%.

## **Introduction:**

The Fashion-MNIST dataset is a popular dataset for benchmarking machine learning algorithms for image classification. It consists of 70,000 grayscale images of 10 different fashion items:

- T-shirt/top
- Trouser
- Pullover
- Dress
- Coat
- Sandal
- Shirt
- Sneaker
- Bag
- Ankle boot.

The images are 28x28 pixels in size and are evenly split between training and test sets.

## Methodology :

The following methodology was used to build the deep learning model:

- ❖ The Fashion-MNIST dataset was loaded into a Python environment.
- ❖ The images were pre-processed by normalising the pixel values to the range [0, 1].
- ❖ A convolutional neural network (CNN) model was created. The CNN model consisted of two convolutional layers, each followed by a max pooling layer.
- ❖ The CNN model was then followed by two fully connected layers.
- ❖ The CNN model was trained on the training set using the Adam optimiser and the categorical cross-entropy loss function.
- ❖ The CNN model was evaluated on the test set.

## Code :

### Step # 1 - Import Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sbn
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Dropout, Dense, Flatten
from keras.optimizers import Adam
from keras.callbacks import TensorBoard
from keras.utils import to_categorical
```

### Step # 2 - Load Data

```
In [2]: fashion_train_df = pd.read_csv('../input/fashion-mnist-datasets/fashion-mnist_train.csv', sep=',')
fashion_test_df = pd.read_csv('../input/fashion-mnist-datasets/fashion-mnist_test.csv', sep=',')
```

```
In [3]: fashion_train_df.shape    # Shape of the dataset
```

```
Out[3]: (60000, 785)
```

```
In [4]: fashion_train_df.columns    # Name of the columns of the DataSet.
```

```
Out[4]: Index(['label', 'pixel1', 'pixel2', 'pixel3', 'pixel4', 'pixel5', 'pixel6',  
              'pixel7', 'pixel8', 'pixel9',  
              ...  
              'pixel775', 'pixel776', 'pixel777', 'pixel778', 'pixel779', 'pixel780',  
              'pixel781', 'pixel782', 'pixel783', 'pixel784'],  
             dtype='object', length=785)
```

```
In [5]: print(set(fashion_train_df['label']))
```

```
{0, 1, 2, 3, 4, 5, 6, 7, 8, 9}
```

```
In [6]: print([fashion_train_df.drop(labels='label', axis=1).min(axis=1).min(),  
              fashion_train_df.drop(labels='label', axis=1).max(axis=1).max()])
```

```
[0, 255]
```

```
In [7]: fashion_train_df.head()
```

```
Out[7]:
```

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	pixel10	pixel11	pixel12	pixel13
0	2	0	0	0	0	0	0	0	0	0	0	0	0	0
1	9	0	0	0	0	0	0	0	0	0	0	0	0	0
2	6	0	0	0	0	0	0	0	5	0	0	0	105	92
3	0	0	0	0	1	2	0	0	0	0	0	114	183	114
4	3	0	0	0	0	0	0	0	0	0	0	0	0	46

```
In [8]: fashion_test_df.shape
```

```
Out[8]: (10000, 785)
```

In [9]:

```
fashion_test_df.head()
```

Out[9]:

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	pixel10	pixel11	pixel12
0	0	0	0	0	0	0	0	0	9	8	0	0	34
1	1	0	0	0	0	0	0	0	0	0	0	0	209
2	2	0	0	0	0	0	0	14	53	99	17	0	0
3	2	0	0	0	0	0	0	0	0	0	161	212	138
4	3	0	0	0	0	0	0	0	0	0	0	37	0

## Step # 3 - Visualisation

In [10]:


```
# Convert the dataframe to numpy array
training = np.asarray(fashion_train_df, dtype='float32')

# Lets show multiple images in a 15x15 grid
height = 10
width = 10

fig, axes = plt.subplots(nrows=width, ncols=height, figsize=(17,17))
axes = axes.ravel() # this flattens the 15x15 matrix into 225
n_train = len(training)

for i in range(0, height*width):
    index = np.random.randint(0, n_train)
    axes[i].imshow(training[index, 1:].reshape(28,28))
    axes[i].set_title(int(training[index, 0]), fontsize=8)
    axes[i].axis('off')

plt.subplots_adjust(hspace=0.5)
```



## Step # 4 - Preprocess Data

```
In [11]: # convert to numpy arrays and reshape
training = np.asarray(fashion_train_df, dtype='float32')
X_train = training[:, 1:].reshape([-1,28,28,1])
X_train = X_train/255 # Normalizing the data
y_train = training[:, 0]

testing = np.asarray(fashion_test_df, dtype='float32')
X_test = testing[:, 1:].reshape([-1,28,28,1])
X_test = X_test/255 # Normalizing the data
y_test = testing[:, 0]
```

```
In [12]: # Split the training set into training and validation sets
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.2, random_state=12345) # TODO : change the random state to 5
```

```
In [13]: # Lets check the shape of all three datasets
print(X_train.shape, X_val.shape, X_test.shape)
print(y_train.shape, y_val.shape, y_test.shape)

(48000, 28, 28, 1) (12000, 28, 28, 1) (10000, 28, 28, 1)
(48000,) (12000,) (10000,)
```

## Step # 5 - Create and Train the Model

### Create the model

```
In [14]: cnn_model = Sequential()
cnn_model.add(Conv2D(filters=64, kernel_size=(3,3), input_shape=(28,28,1), activation='relu'))
cnn_model.add(MaxPooling2D(pool_size = (2,2)))
cnn_model.add(Dropout(rate=0.3))
cnn_model.add(Flatten())
cnn_model.add(Dense(units=32, activation='relu'))
cnn_model.add(Dense(units=10, activation='sigmoid'))
```

### compile the model

```
In [15]: cnn_model.compile(optimizer=Adam(lr=0.001), loss='sparse_categorical_crossentropy', metrics=['accuracy'])
cnn_model.summary()
```

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 26, 26, 64)	640
-----		
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 64)	0
-----		
dropout_1 (Dropout)	(None, 13, 13, 64)	0
-----		
flatten_1 (Flatten)	(None, 10816)	0
-----		
dense_1 (Dense)	(None, 32)	346144
-----		
dense_2 (Dense)	(None, 10)	330
=====		
Total params: 347,114		
Trainable params: 347,114		
Non-trainable params: 0		
-----		

### Train the model

```
In [16]: cnn_model.fit(x=X_train, y=y_train, batch_size=512, epochs=50, validation_data=(X_val, y_val))
```

Train on 48000 samples, validate on 12000 samples

Epoch 1/50  
48000/48000 [=====] - 5s 108us/step - loss: 0.9155 - acc: 0.6463 - val\_loss: 0.4851 - val\_acc: 0.8273

Epoch 2/50  
48000/48000 [=====] - 1s 19us/step - loss: 0.4501 - acc: 0.8417 - val\_loss: 0.4318 - val\_acc: 0.8476

Epoch 3/50  
48000/48000 [=====] - 1s 19us/step - loss: 0.4012 - acc: 0.8586 - val\_loss: 0.3718 - val\_acc: 0.8723

Epoch 4/50  
48000/48000 [=====] - 1s 19us/step - loss: 0.3663 - acc: 0.8720 - val\_loss: 0.3409 - val\_acc: 0.8823

Epoch 5/50  
48000/48000 [=====] - 1s 20us/step - loss: 0.3465 - acc: 0.8787 - val\_loss: 0.3285 - val\_acc: 0.8872

Epoch 6/50  
48000/48000 [=====] - 1s 19us/step - loss: 0.3311 - acc: 0.8838 - val\_loss: 0.3286 - val\_acc: 0.8863

Epoch 7/50  
48000/48000 [=====] - 1s 20us/step - loss: 0.3178 - acc: 0.8874 - val\_loss: 0.3054 - val\_acc: 0.8928

Epoch 8/50  
48000/48000 [=====] - 1s 19us/step - loss: 0.3032 - acc: 0.8938 - val\_loss: 0.3039 - val\_acc: 0.8913

Epoch 9/50  
48000/48000 [=====] - 1s 19us/step - loss: 0.2965 - acc: 0.8956 - val\_loss: 0.2970 - val\_acc: 0.8967

Epoch 10/50  
48000/48000 [=====] - 1s 19us/step - loss: 0.2854 - acc: 0.8985 - val\_loss: 0.2873 - val\_acc: 0.8996

Epoch 11/50  
48000/48000 [=====] - 1s 19us/step - loss: 0.2783 - acc: 0.9025 - val\_loss: 0.2889 - val\_acc: 0.8967

Epoch 12/50  
48000/48000 [=====] - 1s 19us/step - loss: 0.2727 - acc: 0.9029 - val\_loss: 0.2848 - val\_acc: 0.8992

## Step # 6 - Evaluate the Model

```
In [17]: eval_result = cnn_model.evaluate(X_test, y_test)
print("Accuracy : {:.3f}".format(eval_result[1]))
```


10000/10000 [=====] - 1s 55us/step  
Accuracy : 0.918

### Visualize the model's predictions

```
In [18]: y_pred = cnn_model.predict_classes(x=X_test)
```

```
In [19]: height = 10
width = 10


fig, axes = plt.subplots(nrows=width, ncols=height, figsize=(20,20))
axes = axes.ravel()
for i in range(0, height*width):
    index = np.random.randint(len(y_pred))
    axes[i].imshow(X_test[index].reshape((28,28)))
    axes[i].set_title("True Class : {:0.0f}\nPrediction : {:d}".format(y_test[index], y_pred[index]))
    axes[i].axis('off')
plt.subplots_adjust(hspace=0.9, wspace=0.5)
```



## Plot Confusin Matrix

```
In [20]: cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(10,5))
sbn.heatmap(cm, annot=True)
```

```
Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x7feb11cd26d8>
```



### Classification Report

In [21]:

```
num_classes = 10
class_names = ["class {}".format(i) for i in range(num_classes)]
cr = classification_report(y_test, y_pred, target_names=class_names)
print(cr)
```

	precision	recall	f1-score	support
class 0	0.88	0.85	0.86	1000
class 1	0.99	0.99	0.99	1000
class 2	0.91	0.84	0.87	1000
class 3	0.92	0.93	0.92	1000
class 4	0.86	0.91	0.88	1000
class 5	0.98	0.98	0.98	1000
class 6	0.75	0.80	0.78	1000
class 7	0.95	0.96	0.96	1000
class 8	0.98	0.97	0.98	1000
class 9	0.96	0.96	0.96	1000
accuracy			0.92	10000
macro avg	0.92	0.92	0.92	10000
weighted avg	0.92	0.92	0.92	10000

## Conclusion :

The deep learning model was able to achieve an accuracy of 91.8% on the test set. This is a significant improvement over the baseline accuracy of 50%. The model could be further improved by using a larger dataset, a more complex model, or a different optimisation algorithm.

**...THE END...**